

Digital Design and Computer Architecture

A complete overview of the materials taught in the Digital Design and Computer Architecture course of the Basisjahr's second semester at ETH Zürich.

Erik Girdauskas

Table of Contents

Boolean Algebra and Basic Logic	1
Combinational Logic Blocks	4
Sequential Logic	7
Memory Types	9
Finite State Machines	10
Timing, Delay, and Clocking	13
Combinational Timing	14
Sequential Timing	15
Verilog	17
Sequential Logic	19
FSMs, Testbenches, and Simulation	20
Von Neumann Model	21
LC-3 - Expanded Example	25
Data Movement and Addressing Modes	25
Control Instructions and Branching	26
The MIPS ISA	27
MIPS Instruction Formats	28
Basic Instructions	29
Control Instructions and Advanced Concepts	30
Procedures	31
Stack	32
Recursion	32
Microarchitecture	33
Single-Cycle CPU	34
Multicycle CPU	37
Pipelining	39
Data Hazards	41
Branch Hazards	43
Exceptions and Precise State	44
Out-of-Order Execution	45
Reorder Buffer and Register Renaming	46
Dataflow	47
Tomasulo's Algorithm	47
Superscalar Execution	51
Branch Prediction	52
Dynamic Branch Prediction	54
VLIW	56
Systolic Arrays	57

SIMD	58
Array Processors	58
Vector Processors	59
GPUs	62
Memory Hierarchy	66
Caches	69
Virtual Memory	73
Optional x86 Section	82
Additional Exam Problems	84

Color code: major topics, important stuff

by Erik Girdauskas

semester recap - DDCA



foreword: if there is one subject to trust me on, it's this! the sheet is based on my reading notes + the lectures

Basics/Boolean Algebra:

Transistors: The most basic building blocks of computers, they are mostly (only?) implemented with

MOS → metal, oxide (as insulator), semiconductors

→ They are used to build logic gates, don't worry about constructing anything with them for the exam though!



→ 3V is a stand-in for any HIGH voltage here

→ think of n-type transistors as active high and active low components

→ exercise: the 7-segment display in our FPGA displays a value when its transistors are driven to a LOW voltage. What transistors does it use?


Logic Gates:

→ They are basic modules made from transistors and boolean algebra is performed on them

→ the term CMOS stands for complementary MOS, describing the use of both p-type and n-type transistors in larger structures like gates

NOT gate:


(logical inverted)



Truth Table (these all use truth tables):

A	P	N	Y
0	ON	OFF	1
1	OFF	ON	0


general structure of gate schematic →



← example path of 0V being passed as input

NAND gate:


(not and)



TT:

A	B	P1	P2	N1	N2	Y
0	0	ON	ON	OFF	OFF	1
0	1	ON	OFF	OFF	ON	1
1	0	OFF	ON	ON	OFF	1
1	1	OFF	OFF	ON	ON	0

AND gate:




TT:

A	B	Y
0	0	0
0	1	0
1	0	0
1	1	1

→ comment: see how we built this from a NAND gate and an inverter? This saves us transistors compared to building it from scratch

→ NAND and NOR are logically complete because you can build any Boolean expression using those two gates

CMOS gate structure:



→ Example: in XOR, only one network should be on, if both are on, there's a short-circuit
if neither are, an undefined floating output is sent out

→ pMOS only carries 1's (HIGH voltage) well, nMOS only carries 0's (LOW voltage) well

→ this means that with pMOS, passing a logical LOW causes voltage loss (other way around for nMOS)

→ it's also the reason we need both (CMOS)

Electrical stuff (maybe exam relevant?)

→ you can compare running something in series electrically to a logical AND and electrical parallel to a logical OR

→ Latency: Series puts more resistance on the wire carrying it, and as resistance determines latency, series is slower

→ This is why practically, as little series as possible is used

→ Dynamic Power Consumption: power to charge a transistor while the signal changes (inversion): $C \cdot V^2 \cdot f$
Capacitance, Voltage, Charge frequency of capacitor

→ Static Power Consumption: power used while signals do not change: $V \cdot I_{\text{leakage}}$ ← current that keeps flowing, hard to minimize

→ Energy Consumption: Power • Time

→ at no point is the power used meant to exceed the Thermal Design Product (TDP, max. intended power)

Diagram for Gates and their Truth Tables (from now on, we won't use the transistor-level schematic)

Buffer

A	Z
0	0
1	1

AND

A	B	Z
0	0	0
0	1	0
1	0	0
1	1	1

OR

A	B	Z
0	0	0
0	1	1
1	0	1
1	1	1

XOR

A	B	Z
0	0	0
0	1	1
1	0	1
1	1	0

Inverter

A	Z
0	1
1	0

NAND

A	B	Z
0	0	1
0	1	1
1	0	1
1	1	0

NOR

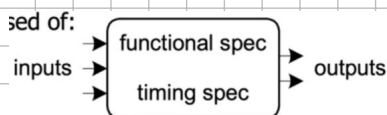
A	B	Z
0	0	1
0	1	0
1	0	0
1	1	0

XNOR

A	B	Z
0	0	1
0	1	0
1	0	0
1	1	1

Circuits:

Structure:



Functional spec: putting outputs in terms of inputs (using the math. def. of function)

→ this uses Boolean algebra, which is mostly repetition from DiskMath

Axioms and laws: (+ replaces v from DM, • replaces ∧, $\bar{A}$ replaces $\neg A$)

All axioms are given in dual form:

□ A dual of a Boolean expression is derived by replacing

- Every AND operation with... an OR operation
- Every OR operation with... an AND
- Every constant 1 with... a constant 0
- Every constant 0 with... a constant 1
- But don't change any of the literals or play with the complements!

Example

$$a \cdot (b + c) = (a \cdot b) + (a \cdot c)$$

$$\rightarrow a + (b \cdot c) = (a + b) \cdot (a + c)$$

(Refers to the entire Boolean expression)

Formal version

1. B contains at least two elements, 0 and 1, such that $0 \neq 1$
2. Closure $a, b \in B$,
(i) $a + b \in B$
(ii) $a \cdot b \in B$
3. Commutative Laws: $a, b \in B$,
(i) $a + b = b + a$
(ii) $a \cdot b = b \cdot a$
4. Identities: $0, 1 \in B$
(i) $a + 0 = a$
(ii) $a \cdot 1 = a$
5. Distributive Laws:
(i) $a + (b \cdot c) = (a + b) \cdot (a + c)$
(ii) $a \cdot (b + c) = a \cdot b + a \cdot c$
6. Complement:
(i) $a + \bar{a} = 1$
(ii) $a \cdot \bar{a} = 0$

Operations with 0 and 1:

1. $X + 0 = X$
2. $X + 1 = 1$

Dual

- 1D. $X \cdot 1 = X$
- 2D. $X \cdot 0 = 0$

AND, OR with identities gives you back the original variable or the identity

Idempotent Law:

3. $X + X = X$

- 3D. $X \cdot X = X$

AND, OR with self = self

Involution Law:

4. $\overline{\overline{X}} = X$

double complement = no complement

Laws of Complementarity:

5. $X + \bar{X} = 1$

- 5D. $X \cdot \bar{X} = 0$

AND, OR with complement gives you an identity

Commutative Law:

6. $X + Y = Y + X$

- 6D. $X \cdot Y = Y \cdot X$

Just an axiom...

Associative Laws:

7. $(X + Y) + Z = X + (Y + Z) = X + Y + Z$

- 7D. $(X \cdot Y) \cdot Z = X \cdot (Y \cdot Z) = X \cdot Y \cdot Z$

Parenthesis order does not matter

Distributive Laws:

8. $X \cdot (Y + Z) = (X \cdot Y) + (X \cdot Z)$

- 8D. $X + (Y \cdot Z) = (X + Y) \cdot (X + Z)$

Axiom

Simplification Theorems:

9. $X \cdot Y + X \cdot \bar{Y} = X$

- 9D. $(X + Y) \cdot (X + \bar{Y}) = X$

10. $X + X \cdot Y = X$

- 10D. $X \cdot (X + Y) = X$

11. $(X + \bar{Y}) \cdot Y = X \cdot Y$

- 11D. $(X \cdot \bar{Y}) + Y = X + Y$

Useful for simplifying expressions

DeMorgan's Law:

$$12. \overline{(X + Y + Z + \dots)} = \bar{X} \cdot \bar{Y} \cdot \bar{Z} \cdot \dots$$

$$12D. \overline{(X \cdot Y \cdot Z \cdot \dots)} = \bar{X} + \bar{Y} + \bar{Z} + \dots$$

Canonical function forms:

Sum of Products (SOP, identical to DNF from DiskMath):

→ Two-level logic: at most two levels of gates → no path goes through more than 2 gates (SOP or Product of Sums (POS))

→ all Boolean equations can be put into SOP and POS

→ deriving it from a truth table: find every entry that evaluates to 1, make these an implicant (connected with +), then make a disjunction of implicants (connect with +)

→ each implicant here is a minterm (product of all input variables)

→ standard notation: enumerate over minterms based on the binary number created by the input pattern

→ Example: $\sum m(0, 6, 7)$

POS: (identical to CNF)

→ deriving it from a truth table: take each expression that evaluates to false, disjoin them (connect with +, this makes them maxterms (sum of all input variables) and conjoin the maxterms (using $\cdot$))

→ standard notation: similar to SOP, but take the product of maxterms (like: $\prod M(0, 6, 7)$)

Conversion rules:

1. **Minterm to Maxterm conversion:**
rewrite minterm shorthand using maxterm shorthand
replace minterm indices with the indices not already used
E.g., $F(A, B, C) = \sum m(3, 4, 5, 6, 7) = \prod M(0, 1, 2)$
2. **Maxterm to Minterm conversion:**
rewrite maxterm shorthand using minterm shorthand
replace maxterm indices with the indices not already used
E.g., $F(A, B, C) = \prod M(0, 1, 2) = \sum m(3, 4, 5, 6, 7)$

3. Expansion of F to expansion of $\bar{F}$:

$$\text{E.g., } F(A, B, C) = \sum m(3, 4, 5, 6, 7) \longrightarrow \bar{F}(A, B, C) = \sum m(0, 1, 2)$$

$$= \prod M(0, 1, 2) \longrightarrow = \prod M(3, 4, 5, 6, 7)$$


4. **Minterm expansion of F to Maxterm expansion of $\bar{F}$:**
rewrite in Maxterm form, using the same indices as F

$$\text{E.g., } F(A, B, C) = \sum m(3, 4, 5, 6, 7) \longrightarrow \bar{F}(A, B, C) = \prod M(3, 4, 5, 6, 7)$$

$$= \prod M(0, 1, 2)$$

■ SOP (sum-of-products) leads to two-level logic

■ Example: $Y = (\bar{A} \cdot \bar{B} \cdot \bar{C}) + (A \cdot \bar{B} \cdot \bar{C}) + (A \cdot \bar{B} \cdot C)$



Combinational Blocks:


→ constructed using gates, they are a layer of abstraction above gates, so gate-level details are hidden

→ Note: Combinational Logic has no memory

Decoder: Has a separate output for every possible input pattern, n inputs, 2^n outputs

2-to-4 decoder:

A_1	A_0	Y_3	Y_2	Y_1	Y_0
0	0	0	0	0	1
0	1	0	0	1	0
1	0	0	1	0	0
1	1	1	0	0	0



[In Verilog:

```
Verilog
module decoder #(parameter N = 3)
  (input [N-1:0] a,
   output reg [2*N-1:0] y);

  always @ (*)
  begin
    y = 0;
    y[a] = 1;
  end
endmodule
```


set output indexed at numerical value of input a to 1

→ oh you will see how it comes up in a CPU


Multiplexer (MUX): core idea: select one of N inputs to mirror to the output, decision is made by the $\log_2 N$ -bit control input:

2-to-1 MUX:

S	D_1	D_0	Y
0	0	0	0
0	0	1	1
0	1	0	0
0	1	1	1
1	0	0	0
1	0	1	0
1	1	0	1
1	1	1	1



(gate-level):



In Verilog (2-to-1 MUX):

```
Verilog
module mux2
  #(parameter width = 8)
  (input [width-1:0] d0, d1,
   input [width-1:0] s,
   output [width-1:0] y);


  assign y = s ? d1 : d0;
endmodule
```

default value
depends on width
we use assign w variable to decide module's width

Simpler TI:


S	C
0	A
1	B

4-to-1 mux:



→ **Lookup Tables (LUT)**, which our FPGAs use prominently for reprogrammable comb. logic, are implemented as trees of MUXes

3-bit input LUT (3-LUT)



3-LUT can implement any 3-bit input function


A function that outputs '1' when there are at least two '1's in a 3-bit input

In C:

```
int count = 0;
for(int i = 0; i < 3; i++) {
  count += input & 1;
  input = input >> 1;
}
if(count > 1) return 1;
return 0;
```

In Hardware (e.g., FPGA):

```
switch(input){
  case 0:
  case 1:
  case 2:
    return 0;
  default:
    return 1;
}
```



Full Adder: our way of performing bitwise addition on two inputs (A,B) with a carry C

gate-level:





Figure 3.15 Gate-level description of a full adder

→ AND gates for each combination of "current-bit" inputs,

OR gates for "current-bit" output S_i and "next-bit" carry C_{i+1}

TT:

a_i	b_i	carry _i	carry _{i+1}	S_i
0	0	0	0	0
0	0	1	0	1
0	1	0	0	1
0	1	1	1	0
1	0	0	1	1
1	0	1	1	0
1	1	0	1	0
1	1	1	1	1



(Optional explanation)

For multi-bit addition, the most efficient CL (comb. logic) block is the **Carry Lookahead Adder**:




Figure 5.5 32-bit ripple-carry adder

↳ in M&H Chapter 5, you see the **Ripple-Carry Adder**, which has more latency as the addition is done sequentially from the MSB (most significant bit) to the LSB (least significant bit), so the delay Tripple grows with the number of bits (linearly)

The **Carry Lookahead Adder (CLA)** eliminates this delay:

→ the adder is **divided** into blocks

→ circuitry can quickly determine C_{out} of a block with known C_{in}

→ **Generate (G)** and **Propagate (P)** signals describe how C_{out} is determined

→ the i -th column of the CLA generates a carry-out if it produces one independent of C_{in} → C_i is generated if $A_i = 1$ and $B_i = 1$ ⇒ $A_i B_i$

→ it **propagates** a carry-out if $A_i = 1$ or $B_i = 1$ ⇒ $A_i + B_i$

→ the i -th column generates C_i if it generates a carry or propagates a carry in ($P_i C_{i-1}$) ⇒ $C_i = A_i B_i + (A_i + B_i) C_{i-1}$
 $= G_i + P_i C_{i-1}$

→ our definitions from now extend to **multi-bit blocks**: a block generates a carry if it produces one independent of C_{in} and propagates one if a carry is produced under the condition of $C_{in} = 1$

→ G and P extend to $G_{i:j}$ and $P_{i:j}$ for the block $i:j$

→ **pattern**: block generates a carry if most sig. column generates one or the most sig. column generates one and the next sig. column generates one

$$G_{3:0} = G_3 + P_3(G_2 + P_2(G_1 + P_1 G_0))$$



→ it **propagates** one if all columns propagate: $P_{3:0} = P_3 P_2 P_1 P_0$

$$C_{i:j} = G_{i:j} + P_{i:j} C_j$$

→ single-bit and block generate/propagate is computed simultaneously

→ **Critical path** (defined later): Compute $G_{3:0}$ and G_0 in the first block, then advance sequentially

→ **Timing definition** (also formally introduced later): $t_{CLA} = t_{pg} + t_{pg_block} + (\frac{N}{k} - 1) \cdot t_{AND_OR} + k \cdot t_{FA} \leftarrow k \text{ full adders}$
 (for N bits in k-bit blocks)
 individual P,G gates (single AND/OR construct)



Verilog

```
module adder #(parameter N = 8)
    (input [N-1:0] a, b,
     input cin,
     output [N-1:0] s,
     output cout);

    assign {cout, s} = a + b + cin;
endmodule
```

↑ bundled parameters → $\{a, b\}$

Verilog adder

I will just continuously leave these on the notes to comb. blocks

Programmable Logic Array (PLA):

The diagram shows a logic circuit with three inputs labeled A, B, and C. These inputs are connected to eight AND gates arranged in a vertical column. Each AND gate has two inputs, and the connections between the inputs and the gates are shown as a grid of lines. The outputs of the eight AND gates are connected to a central gray rectangular block labeled "Connections". From the right side of this block, three lines emerge, each connected to an OR gate. The outputs of these OR gates are labeled X, Y, and Z.

OR gates $\rightarrow$ amount of output columns (we call it M) ($\leq 2^n$)

→ you can base this on truth tables



The diagram shows a logic circuit with 10 AND gates and 3 OR gates. On the left, three inputs labeled A, B, and C are connected to the first three AND gates. The remaining seven AND gates have various combinations of these inputs connected to them. The outputs of all 10 AND gates are connected to a central grey block labeled 'Connections'. From this block, three lines lead to three OR gates on the right, which are labeled X, Y, and Z.

We do not need this output

a_i	b_i	$carry_i$	$carry_{i+1}$	S_i
0	0	0	0	0
0	0	1	0	1
0	1	0	0	1
0	1	1	1	0
1	0	0	0	1
1	0	1	1	0
1	1	0	1	0
1	1	1	1	1



→ {NAND, NOR} are also log. complete

Comparator:



Our first lab!

Magnitude comparator:



N-bit magnitude comparator in Verilog:

synthesizes to

This style is known as behavioral Verilog, we will go into it in a bit

ALU: (Arithmetic Logic Unit)

```
module comparators # (parameter N = 8)
    (input [N-1:0] a, b,
    output eq, neq,
    output lt, lte,
    output gt, gte);
```

```
assign eq = (a == b);
assign neq = (a != b);
assign lt = (a < b);
assign lte = (a <= b);
assign gt = (a > b);
assign gte = (a >= b);
endmodule
```

→ this is crucial to CPUs




Figure 5.14 ALU symbol

Table 5.1 ALU operations

$F_{2:0}$	Function
000	A AND B
001	A OR B
010	A + B
011	not used
100	A AND $\bar{B}$
101	A OR $\bar{B}$
110	A - B
111	SLT


→ operations vary depending on your microarchitecture

Your lab. one had XOR and NOR

aside: 2's complement $\rightarrow$ negative binary numbers start with 1,
zero is always a N-bit zero (5)


→ this is how ALU's do math usually: convert a positive like 0101 to negative → invert bits, then add 1 → 1011 (-5)

Implementation:



Tri-State Buffer:

Tristate Buffer




E	A	Y
0	0	Z
0	1	Z
1	0	0
1	1	1

Figure 2.40 Tristate buffer

→ z is a floating value that is neither 1 nor 0 (more later)

→ example usage: a shared bus → beautiful diagram:



Simplify TT:

E	Y
0	Z
1	A


→ when CPU in use, we set GetMem to 0, it pushes z to the bus, which doesn't interfere with the CPU value

cool usage: MUX using tri-state buffers:

Logic Simplification:


→ why? we want to reduce the number of gates/inputs to save space/latency

→ **Uniting Theorem:** find "ON" set of function, if an input can change without changing the output, its value is not needed



$$Y = D_0 \bar{S} + D_1 S$$

Figure 2.56 Multiplexer using tristate buffers



Key Tool: The Uniting Theorem — $F = \bar{A}\bar{B} + AB$

A	B	F
0	0	0
0	1	0
1	0	1
1	1	1

$$F = \bar{A}\bar{B} + AB = A(\bar{B} + B) = A(1) = A$$

B's value changes within the rows where $F=1$ ("ON set")

A's value does NOT change within the ON-set rows

If an input (B) can change without changing the output, that input value is not needed

→ B is eliminated, A remains

→ B is eliminated, A remains

A	B	G
0	0	1
0	1	0
1	0	1
1	1	0


$$G = \bar{A}\bar{B} + AB = (\bar{A} + A)\bar{B} = \bar{B}$$

B's value stays the same within the ON-set rows

A's value changes within the ON-set rows

→ A is eliminated, B remains

Priority Circuit: Inputs: requestors, Output: who gets priority:




A ₃	A ₂	A ₁	A ₀	Y ₃	Y ₂	Y ₁	Y ₀
0	0	0	0	0	0	0	0
0	0	0	1	0	0	0	1
0	0	1	0	0	0	1	0
0	0	1	1	0	0	1	0
0	1	0	0	0	1	0	0
0	1	0	1	0	1	0	0
0	1	1	0	0	1	0	0
0	1	1	1	0	1	0	0
1	0	0	0	1	0	0	0
1	0	0	1	1	0	0	0
1	0	1	0	1	0	0	0
1	0	1	1	1	0	0	0
1	1	0	0	1	0	0	0
1	1	0	1	1	0	0	0
1	1	1	0	1	0	0	0
1	1	1	1	1	0	0	0

→ simplify:

A ₃	A ₂	A ₁	A ₀	Y ₃	Y ₂	Y ₁	Y ₀
0	0	0	0	0	0	0	0
0	0	0	1	0	0	0	1
0	0	1	X	0	0	1	0
0	1	X	X	0	1	0	0
1	X	X	X	1	0	0	0


Figure 2.29 Priority circuit truth table with don't cares (X's)

circuit:



Sequential Logic:

Cross-Coupled Inverter:




→ Q and $\bar{Q}$ depend on each other

Case 1: $Q=0$; I₂ receives false → $\bar{Q}$ turns true, I₁ receives false → Q false ...

Case 2 analogous → the circuit has 2 stable states (bistable)

→ Metastable state: both outputs halfway between 0 and 1 (explained later)

→ we have no control mechanism, so this concept is expanded into the R-S latch:



Input		Output
R	S	Q
1	1	Q_{prev}
1	0	1
0	1	0
0	0	Forbidden

→ instead of inverters, we cross-coupled NAND gates (book uses NOR)

→ Data stored at Q, its inverse at $\bar{Q}$

→ R and S are control inputs (quiescent/idle) state: both are 1)

→ S(set): you set by driving S to 0 and not resetting (keep R at 1)

(block diagram: $\begin{matrix} R & Q \\ S & \bar{Q} \end{matrix}$)

→ R(reset): you reset by driving R to 0 and not setting (keep S at 1)

→ S and R should not both be at 0. why?

→ logically: you cannot both set and reset at once, this also causes both Q and Q' to be 1, breaking the invariant (here: meaning of Q)

→ practically: if both Q and Q' are set to 1 at the same time, Q and Q' start to oscillate (metastability, values depend on each other)

→ they would just start wildly switching between 1 and 0 because Q cannot equal Q'


→ the time this takes to resolve varies (seen later)

→ The gated D latch guarantees correct operation:

→ book version uses a clock instead of WE

(block diagram: )


→ D is our only real input (data), if WE is enabled, Q takes the value of D




Input		Output
WE	D	Q
0	0	Q _{prev}
0	1	Q _{prev}
1	0	0
1	1	1

→ Verilog latch →

```
Verilog
module latch (input clk,
               input [3:0] d,
               output reg [3:0] q);
    always @ (clk, d)
        if (clk) q <= d;
endmodule
```



D flip-flop: (I'll use the clock version, you can imagine CLK as WE for now)



→ The reason we use this problematic terminology is for the following analogy: "Slave copies Master"

→ L1 and L2 are D latches, L1 is the master, L2 is the slave

→ CLK = 0 → L1 is transparent (taking in the value from D), L2 opaque (keeping its value)


Block diagram:



→ CLK = 1 → L1 is opaque, L2 is transparent (N1 is cut off from D and takes the value stored in Q)

→ in action: L1 copies to L2 at the rising edge (CLK goes from 0 → 1) and remembers its state otherwise

N-bit register: bank of N flip-flops




→ they share CLK and all flip-flops are updated at the rising edge

Adjacent concepts:

Enabled flip-flop → WE decides if Q is written on the rising edge

Resettable flip-flop → RESET input forces Q = 0
→ active low input → something happens after it is at 0

Memory arrays: shared properties of all memories:



→ memory is organized as a 2D array of memory cells

→ a program reads/writes the contents of an array row

→ rows are addressable → an array with N-bit addresses and M-bit data

has 2^N rows and M columns

→ each row is called a word → there are 2^N M-bit words

→ depth → #rows, width → #columns (word size), size → depth · width

→ memory arrays are built as arrays from bit cells

→ each bit cell stores 1 bit of data and is connected to both the wordline and bitline

→ for each combination of address bits, the memory asserts a single wordline that activates all bit cells in the row

→ wordline is HIGH → stored bit is transferred to or from the bitline

→ else, the bitline is not connected to the bit cell

→ usually, the bitline is at $\bar{Z}$ to read, when the wordline is ON, it is set to ON → this connects bitline to stored bit

→ this causes the bitline to overpower and rewrite the bit cell's data

Memory Arrays:


→ read: wordline asserted, bit cells then drive their lines H/L

→ write: drive bitlines H/L, then assert wordline to store the new values

→ ports: all memories have ≥ 1 port for read and write access each



→ a multiported memory allows simultaneous reading/writing

The lectures used a slightly different design:



→ here, the point is to set individual cells rather than rows

→ decoder to decide row, MUX to decide current row → select input



Memory types: (NOTE: this is part 1, part 2 will come around W11-W12 → caches, memory banking...)

DRAM: a bit is stored as the presence or absence of charge in a capacitor

→ nMOS transistor behaves as switch between capacitor and bitline

→ wordline asserted → nMOS turns ON → bit value put on/taken from bitline

→ it is not driven high/low by a transistor tied to V_{DD} or GND (dynamic RAM)
voltage needed for HIGH ground, voltage for LOW

→ on read: data value driven from capacitor to bitline, write: opposite

→ charge leaks over time → DRAM must be reset every few ns when not read

SRAM: static, does not need to be refreshed

→ data bit stored in cross-coupled inverter with outputs bitline, $\overline{\text{bitline}}$ (as Q, $\bar{Q}$)

→ wordline asserted → both nMOS turn on, values are either set or gotten

→ if noise degrades charge, cross-coupled inverters restore values

→ latency and throughput depend on array size



5.5 Memory Arrays




Figure 5.44 DRAM bit cell




Figure 5.45 DRAM stored values




Figure 5.46 SRAM bit cell

Memory Type	Transistors per Bit Cell	Latency
flip-flop	~20	fast
SRAM	6	medium
DRAM	1	slow

important, DRAM is deepest and slowest


[illegible]

→ you will see that caches are basically **fully SRAM**

ROM: a bit is the presence or absence of a transistor

→ reading a cell: when the wordline of a row is ON, pull the bitline HIGH (weakly)

→ If a transistor exists, the corresponding bitline is driven LOW (absence means binary 0)



→ a ROM bit cell is a combinational circuit (no state to remember)

→ any ROM can be constructed from a decoder and NOT gates

→ ROM can perform any two-level logic function (→ recap: level:

→ traditional ROM is hard coded in manufacturing

→ if the transistors are always there, but disconnectable to ground, we have Programmable ROM (PROM)

→ if instead of transistors, we use **floating gates**, we have **(Electrically) Erasable PROM** (EEPROM/EPROM depending on technical details)

In Verilog:

```


module can # (parameter N = 6, M = 32);
    input          clk;
    input          we;
    input [N-1:0]  adr;
    input [M-1:0]  din;
    output [M-1:0] dout;

    reg [M-1:0] mem [2^N-1:0];

    always @ (posedge clk)
        if (we) mem[adr] <= din;

    assign dout = mem[adr];
endmodule

```



```

module rom(input [1:0] adr,
           output reg [2:0] dout);

    always @ (adr)
        case (adr)
            2'b00: dout <= 3'b011;
            2'b01: dout <= 3'b110;
            2'b10: dout <= 3'b100;
            2'b11: dout <= 3'b010;
        endcase
endmodule

```

Back to sequential circuits:


State: It remembers past events and contributes to the output depending on current inputs.

Analogy: Combinational lock:

State A: lock not open, 0/3 inputs
 B: locked, 1/3 complete
 C: locked, 2/3 complete
 D: unlocked

output
 ↓
 ↓ need to be completed in order,
 ↓ else return to A
 ↓

→ state diagram for combination 213, L22, 123
(it's one of the locks you turn left/right)



Asynchronous vs. synchronous:

→ **async.**: immediate state transition after input, state transitions are synchronized by nothing

→ sync.: state transitions at fixed times

→ done in most modern computers, it is much simpler to design, but input is delayed (usually by picoseconds today?)


→ in our course, we always use a clock to synchronize (you've seen this in the memory chapter already)

Clock:

CLK: 1 0

6 ns 3 ns 3 ns 10


→ state only changes at the rising and falling edges in a clock-controlled circuit



→ we use the clock to trigger transitions across many elements in sync. circuit

→ during a clock cycle, combinational logic blocks compute the next state (also known as next-state logic)

→ in the timing chapter, you will see how important it is for a clock cycle to be longer than any combinational delay (computing)
reason why we use flip-flops and not latches in sequential circuits:




(this is for a D latch)

FSM: Discrete time-model of a stateful system

It contains:

- a finite number of states
- a finite number of external inputs
- a finite number of external outputs
- a specification of all state transitions
- a specification of how external outputs are determined



→ The state register(s) store the current state and output the next state at the rising edge (only sequential circuit)

→ **Combinational circuits:** Next state logic and output logic

Moore machines: outputs depend on current state only

Meady machines: outputs also depend on current inputs

FSM design (crucial for exam!):

1. Identify your inputs/outputs

2. Draw your state transition diagram

→ start with the reset state and what is caused by it

→ add states and transitions

→ this is similar to programming

→ FSMs have a sequential control flow (conditionals, jumps)

→ if/else construct controlled by inputs → outputs controlled by state or inputs

3. **Moore machine:** write state transition table, write output table


Meady machine: write state transition and output table (combined)

4. Select state encodings:

Binary (full) encoding: Use minimum possible number of bits → $\log_2(\# \text{states})$ bits
→ Example: 00, 01, 10, 11
→ minimizes # flip-flops, not necessarily output logic

One-hot encoding: one bit for every state → #states bits needed → example: 0001, 0010, 0100, 1000
→ simple, automatable, doesn't need much next state logic, but inefficient for flip-flops

Output encoding: directly encode outputs in state



→ example: 001, 010, 100, 110 for a traffic light (Only works for Moore type)

5. Write Boolean equations for next state and output logic

6. Sketch the circuit (this is formal, you won't have to do this)

Simple example:

■ "Smart" traffic light controller

□ 2 inputs:


■ Traffic sensors: T_A, T_B (TRUE when there's traffic)

□ 2 outputs:

■ Lights: L_A, L_B (Red, Yellow, Green)


□ State changes every 5 seconds

■ Except if green and traffic, stay green



■ Inputs: CLK, Reset, T_A, T_B



■ Outputs: L_A, L_B



■ Moore FSM: outputs labeled in each state


□ States: Circles

□ Transitions: Arcs



Current State		Inputs		Next State	
S_1	S_0	T_A	T_B	S'_1	S'_0
0	0	0	X	0	1
0	0	1	X	0	0
0	1	X	X	1	0
1	0	X	0	1	1
1	0	X	1	1	0
1	1	X	X	0	0


State	Encoding
S0	00
S1	01
S2	10
S3	11



Current State		Inputs		Next State	
S_1	S_0	T_A	T_B	S'_1	S'_0
0	0	0	X	0	1
0	0	1	X	0	0
0	1	X	X	1	0
1	0	X	0	1	1
1	0	X	1	1	0
1	1	X	X	0	0

State	Encoding
S0	00
S1	01
S2	10
S3	11

$$S'_1 = (\overline{S_1} \cdot S_0) + (S_1 \cdot \overline{S_0} \cdot \overline{T_B}) + (S_1 \cdot \overline{S_0} \cdot T_B)$$




Current State		Inputs		Next State	
S_1	S_0	T_A	T_B	S'_1	S'_0
0	0	0	X	0	1
0	0	1	X	0	0
0	1	X	X	1	0
1	0	X	0	1	1
1	0	X	1	1	0
1	1	X	X	0	0

State	Encoding
S0	00
S1	01
S2	10
S3	11


$$S'_1 = S_1 \text{ xor } S_0 \quad (\text{Simplified})$$

$$S'_0 = (\overline{S_1} \cdot \overline{S_0} \cdot \overline{T_A}) + (S_1 \cdot \overline{S_0} \cdot \overline{T_B})$$



Current State		Outputs	
S_1	S_0	L_A	L_B
0	0	green	red
0	1	yellow	red
1	0	red	green
1	1	red	yellow

Output	Encoding
green	00
yellow	01
red	10



Current State		Outputs			
S_1	S_0	L_{A1}	L_{A0}	L_{B1}	L_{B0}
0	0	0	0	1	0
0	1	0	1	1	0
1	0	1	0	0	0
1	1	1	0	0	1

Output	Encoding
green	00
yellow	01
red	10


$$L_{A1} = S_1$$

$$L_{A0} = \overline{S_1} \cdot S_0$$

$$L_{B1} = \overline{S_1}$$

$$L_{B0} = S_1 \cdot S_0$$

Schematic:




Practical Exam Example:

Note: I see it as very unlikely that you will have to draw up a schematic, everything else is important!

from FS22:

(a) [25 points] Complete the Moore-type FSM below by (1) drawing the transition edges between the states (including reset), (2) specifying the edges' respective input bits, and (3) specifying the output bits of each state. Any input for which no outgoing edge is specified is assumed as a self loop. The 6 given states are sufficient, do not draw additional states. Note: Passengers sometimes slip in or out through incorrect doors, or even through closed doors. Your FSM must correctly handle such cases.



outputs: idle 0 $\rightarrow$ 10, idle 1 $\rightarrow$ 10, transit $\rightarrow$ 00, unload 1 $\rightarrow$ 01, unload 2 $\rightarrow$ 01, emergency $\rightarrow$ 11


Timing:

Electrical Basics:

lowest voltage in a system $\rightarrow$ ground/GND (usually 0V)

highest voltage in a system $\rightarrow$ V_{DD} , comes from PSU

Logic Levels:



$\rightarrow$ first gate: driver, second gate: receiver

$\rightarrow$ driver produces logical LOW (0) $\rightarrow$ output between 0 and V_{OL} (output low)

HIGH (1) $\rightarrow$ output between V_{OH} and V_{DD} (output high)

$\rightarrow$ receiver gets input between 0 and V_{IL} $\rightarrow$ reads logical LOW (input low)
 V_{IH} and V_{DD} $\rightarrow$ reads logical HIGH (input high)

$\rightarrow$ forbidden zone: between V_{IL} and V_{IH} (caused by defects or noise)

$\rightarrow$ Noise Margins: $NM_L = V_{IL} - V_{OL}$ (low margin), $NM_H = V_{OH} - V_{IH}$ (high margin). Reasoning: we need output high $>$ input high and output low $<$ input low to correctly interpret outputs (with noise).

before the clock edge and t_{hold} after $\Rightarrow$ aperture time = $t_{setup} + t_{hold}$

Dynamic Discipline: the inputs of a sync. circuit must be stable during the aperture time

$\rightarrow$ sample inputs that are not changing

System Timing: clock period T_c = time between rising edges of a clock, $1/T_c = f_c \rightarrow$ clock frequency, which is measured in Hz (cycles per second)

$\rightarrow$ in this image, we try to calculate the clock period

$\rightarrow R_1$ produces Q_1 on the rising edge, they go through a block of comb. logic that produces D_2 , which goes to R_2

$\rightarrow$ gray arrows $\rightarrow$ contamination delay, blue arrows $\rightarrow$ propagation delay

Setup Time Constraint:

$\rightarrow D_2$ must be settled no later than before t_{setup}

$\rightarrow$ we know that $T_c \geq t_{pcq} + t_{pd} + t_{setup} \rightarrow$ can be rearranged to $t_{pd} \leq T_c - (t_{pcq} + t_{setup})$

this equation is called setup time constraint,

because the comb. logic block cannot meaningfully compute D_2 during t_{pcq} and t_{setup}

$\rightarrow$ if t_{pd} is too long, D_2 may not be at its final value when needed $\rightarrow$ malfunction $\rightarrow$ solution: increase T_c or redesign CL

Hold Time Constraints:

$\rightarrow$ the designer usually doesn't have control over f_{ccq} , t_{pcq} , t_{hold} , t_{setup} , or T_c depends on manufacturers dictated internally

$\rightarrow$ hold time constraint: D_2 must not change for t_{hold} after rising edge, but it

can change $f_{ccq} + t_{cd}$ after rising edge $\Rightarrow t_{ccq} + t_{cd} \geq t_{hold}$

$\Rightarrow t_{cd} \geq t_{hold} - t_{ccq}$

t_{ccq} and t_{hold} outside our control

Example: 2 flip-flops: no comb. logic between $\rightarrow t_{hold} \leq t_{ccq}$, in practice, $t_{hold} = 0$ (or any chain of seq. logic elements)


$\rightarrow$ hold time violations cannot be fixed without redesigning the circuit (by increasing contamination delay) $\rightarrow$

cannot be fixed by adjusting clock period $\rightarrow$ very serious for real circuits

Clock skew (this came up in the lecture, so we must examine it)


$\rightarrow$ clock edges may vary (clock skew), caused by differing clock-register wire length, noise, and using both gated and ungated clocks

$\rightarrow$ here, CLK_2 is early because of the routing of the clock wire between R_1 and R_2




$\rightarrow$ clock may arrive up to t_{skew} earlier

$\rightarrow$ effect on setup time constraint:




→ worst case: R_1 receives latest skewed clock, R_2 receives earliest skewed clock because comb. logic propagation has least time

→ data still has to setup and propagate before sampling $\Rightarrow T_c \geq t_{pcq} + t_{pd} + t_{setup} + t_{skew}$
 $\Rightarrow t_{pd} \leq T_c - (t_{pcq} + t_{setup} + t_{skew}) \rightarrow$ skew added to constraint



Effect on hold time constraint:



→ worst case: R_1 gets earliest skewed clock, R_2 receives latest skewed clock

$\Rightarrow$ data must not arrive until hold time after late clock

$\Rightarrow t_{pcq} + t_{cd} \geq t_{hold} + t_{skew} \Rightarrow t_{cd} \geq t_{hold} + t_{skew} - t_{pcq}$

→ overall, clock skew increases both setup and hold time

Metastability: result of violating ^(in flip-flops) dyn. discipline $\Rightarrow$ output can take on voltage between 0 and V_{DD} in forbidden zone $\rightarrow$ metastable state

→ flip-flop will resolve output to 1 or 0, but resolution time is unbounded

Resolution time: when a flip-flop input changes at random time during clock cycle, t_{res} (resolution time is a random variable) (Aww flashbacks! 🤔)

→ input changes outside aperture $\rightarrow t_{res} = t_{pcq}$, input changes within aperture $\rightarrow t_{res}$ can be longer

Equation for t substantially longer than $t_{pcq} \rightarrow P(t_{res} > t) = \frac{T_0}{T_c} e^{-\frac{t}{T_0}}$ → T_0 and τ are characteristic to flip-flop

→ $\frac{T_0}{T_c} \rightarrow$ prob. that input changes at a bad time

Synchronizers:


→ asynchronous inputs like human inputs are inevitable in larger systems $\rightarrow$ needs handling to prevent metastability

→ a synchronizer receives async. input D and CLK, produces Q in bounded time, the output has a valid logic level with near certainty (if D stable, $Q = D$, else choose HIGH or LOW)

→ in the image, F_1 samples D on the rising edge of CLK (during aperture, momentarily metastable)
 → Over the clock period, D_2 resolves, F_2 samples D_2 , we get a good output Q

→ synchronizer fails when Q becomes metastable $\rightarrow t_{res} > T_c - t_{setup}$

→ $P(\text{failure}) = \frac{T_0}{T_c} e^{-\frac{T_c - t_{setup}}{T_0}}$, $P(\text{failure})/\text{sec} = N \frac{T_0}{T_c} e^{-\frac{T_c - t_{setup}}{T_0}}$
 ↑ times D changes per second
 (MTBF)



synchronizer circuit

→ mean time between failures usually measures system reliability $\Rightarrow \text{MTBF} = \frac{1}{P(\text{failure}/\text{sec})} = \frac{T_c e^{\frac{T_c - t_{setup}}{T_0}}}{N T_0}$

Some notes on parallelism (pipelining will be introduced)

Token $\rightarrow$ group of inputs processed to produce group of outputs

Latency $\rightarrow$ time for 1 token to pass from start to end

Throughput $\rightarrow$ number of tokens per unit time


Spatial parallelism: multiple copies of hardware, Temporal parallelism (pipelining): task broken into stages, multiple tasks spread across stages $\rightarrow$ different task in each stage at once

$\rightarrow$ assume a latency L for a nonparallel task with a throughput of $1/L$. With N copies of HW, it becomes N/L , which is the ideal case for temporal par. too, but in practice, you get $1/L_1 \leftarrow$ longest latency for a step

$\rightarrow$ Exam practicality: Prof. Motlu hasn't made a timing task yet, but they do come up in older exams

Nice trial question: (PS15), we're looking at e)

(d) (4 points) Draw a gate-level schematic that realizes the function Z using only 2-input AND, OR gates. Assume that you have all the inputs (A, B, C, D) and their complements ($\bar{A}, \bar{B}, \bar{C}, \bar{D}$) available.



(e) (2 points) Assume that all the gates (AND, OR, NOT) in the previous diagram have a propagation delay of 100 ps and a contamination delay of 50 ps. What is the delay of the longest (critical) path and the shortest path of this circuit?

$\rightarrow$ in general, I would focus on the basics, less on skew/synchronizers

next section: Verilog!

$\rightarrow$ This may be asked: Verilog is not a programming language, it is a HDL (hardware description language) to allow for synthesis of digital circuits, but it often features similar syntax

$\rightarrow$ a module is a HW block that needs defined inputs/outputs

$\rightarrow$ Behavioral Verilog: using logical operators to describe what a module does

$\rightarrow$ Structural Verilog: gate-level description at the bottom level, at a higher level it is about building modules using instances of modules

$\rightarrow$ you start with name and listing I/O

$\rightarrow$ the assign LHS = RHS statement is used for CL outside always statements


$\rightarrow \neg \rightarrow$ NOT, $\& \rightarrow$ AND, $| \rightarrow$ OR, $\sim \& \rightarrow$ NAND, $\sim | \rightarrow$ NOR, $\wedge \rightarrow$ XOR, $\sim \wedge \rightarrow$ XNOR

$\rightarrow$ Verilog signals (individual bits of variables) can be set to 0, 1, z , x
floating unknown

$\rightarrow$ Simulation: checking modules for logical correctness

$\rightarrow$ Synthesis: HDL code is transformed to a netlist describing hardware produced

$\rightarrow$ code may be simplified during synthesis, there is a schematic output like so:



$\rightarrow$ a top module is the main piece of hardware, can consist of

CL blocks, registers, FSMs

$\rightarrow$ CL: bitwise operators act on single-bit signals or multi-bit busses

module inv (input [3:0] a, output [3:0] y);
assign y = ~a;
endmodule

$a[3:0] \leftarrow$ 4-bit bus, given in little-endian order ($a[3]$ is MSB, $a[0]$ is LSB)
 $a[0:3] \leftarrow$ big-endian, also permitted
 $\rightarrow$ just be consistent



```
module sillyfunction (input a, b, c,
                      output y);
    assign y = ~a & ~b & ~c |
               a & ~b & ~c |
               a & ~b & c;
endmodule
```

Verilog Operator	Name	Functional Group
[]	bit-select or part-select	
()	parenthesis	
!	logical negation	logical
~	negation	bit-wise
&	reduction AND	reduction
	reduction OR	reduction
~&	reduction NAND	reduction
~	reduction NOR	reduction
^	reduction XOR	reduction
~^ or ^~	reduction XNOR	reduction
+	unary (sign) plus	arithmetic
-	unary (sign) minus	arithmetic
{ }	concatenation	concatenation
{ { } }	replication	replication
*	multiply	arithmetic
/	divide	arithmetic
%	modulus	arithmetic
+	binary plus	arithmetic
-	binary minus	arithmetic
<<	shift left	shift
>>	shift right	shift
>	greater than	relational
>=	greater than or equal to	relational
<	less than	relational
<=	less than or equal to	relational
==	case equality	equality
!=	case inequality	equality
&	bit-wise AND	bit-wise
^	bit-wise XOR	bit-wise
	bit-wise OR	bit-wise
&&	logical AND	logical
	logical OR	logical
?:	conditional	conditional


```

module gates (input [3:0] a, b,
              output [3:0] y1, y2,
              y3, y4, y5);

  /* Five different two-input logic
  gates acting on 4 bit busses */
  assign y1 = a & b; // AND
  assign y2 = a | b; // OR
  assign y3 = a ^ b; // XOR
  assign y4 = ~(a & b); // NAND
  assign y5 = ~(a | b); // NOR
endmodule

```

→ you know the code, the assign statements we saw are continuous, end with ;

→ when inputs on RHS change, LHS recomputed

→ Verilog comments are exactly like Java

→ white space is irrelevant, module names can't begin with numbers, vars. are case-sensitive

reduction operators → using the operators we just introduced in a unary way

conditional assignment: you know this as the ternary operator

```

module mux4 (input [3:0] d0, d1, d2, d3,
             input [1:0] s,
             output [3:0] y);

  assign y = s[1] ? (s[0] ? d3 : d2)
              : (s[0] ? d1 : d0);
endmodule

```

i did this for Lab 5

```

module and8 (input [7:0] a,
             output y);

  assign y = &a;

  // &a is much easier to write than
  // assign y = a[7] & a[6] & a[5] & a[4] &
  // a[3] & a[2] & a[1] & a[0];
endmodule

```

```

module fulladder (input a, b, cin,
                 output s, cout);

  wire p, g;

  assign p = a ^ b;
  assign g = a & b;

  assign s = p ^ cin;
  assign cout = g | (p & cin);
endmodule

```

boolean equations for this full adder: $P = A \oplus B$, $G = A \cdot B$, $S = P \oplus C_{in}$, $C_{out} = G + P \cdot C_{in}$

→ the continuous assign statements are concurrent → For CL, order doesn't matter

→ wires are declared to hold internal vars. like P and G, they only need to be declared for

multi-bit busses but are good practice

Numbers: (the book won't tell you this: variables can have the data types integer and real → Example: real seven = 7.00;)

Generally: number format: N's value, size in bits, bases: 'b' → binary, 'o' → octal, 'd' → decimal, 'h' → hex

→ size can be omitted, then a number is assumed to have as many bits as its expression

Table 4.5 Verilog AND gate truth table with z and x

A	z			
	0	1	x	x
0	0	0	0	0
1	0	1	x	x
x	0	x	x	x

Bit Swizzling:

assign y = {2:1, {3{d[0]}}, c[0], 3'b10};

→ { 3 concatenates busses and bits → example: {3{d[0]}} copies d[0] thrice (and constants)

Delays: ignored in synthesis, used in simulation and testing → ensure timing is correct

→ of course, electrical factors in gate delays cannot be considered

```

`timescale 1ns/1ps

module example (input a, b, c,
                output y);
  wire ab, bb, cb, n1, n2, n3;

  assign #1 ab = a & b;
  assign #2 n1 = ab & bb & cb;
  assign #2 n2 = a & bb & cb;
  assign #2 n3 = a & bb & c;
  assign #4 y = n1 | n2 | n3;
endmodule

```

→ timescale is necessary in a file i believe, format: 'timescale unit/precision'

→ here: every unit is 1ns, sim. has 1ps precision

→ # indicates number of units of delay → here: and → 2ns delay, or → 1ns delay

→ can be placed in assign, non-blocking, blocking assignments

structural modeling:

→ the book actually employs bad practice for instantiation:

```

module mux2 (input [3:0] d0, d1,
             input s,
             output [3:0] y);
  tristate t0 (d0, s, y);
  tristate t1 (d1, s, y);
endmodule

module mux4 (input [3:0] d0, d1, d2, d3,
             input [1:0] s,
             output [3:0] y);
  wire [3:0] low, high;
  mux2 lowmux (d0, d1, s[0], low);
  mux2 highmux (d2, d3, s[0], high);
  mux2 finalmux (low, high, s[1], y);
endmodule

```

you can optimize in instantiations

3 instances of mux2

→ good practice: submodule name (.input1(i1), .output1(o1))
example → mux2 lowmux(.d0(d0), .d1(d1), .s(s[0]), .y(low));

→ accessing parts of busses: use array notation following your endian format → d0[7:4]

(example)

```

module mux2_8 (input [7:0] d0, d1,
               input s,
               output [7:0] y);
  mux2 lsbmux (d0[3:0], d1[3:0], s, y[3:0]);
  mux2 msbmux (d0[7:4], d1[7:4], s, y[7:4]);
endmodule

```


Sequential blocks:

```
module flop (input clk,
             input [3:0] d,
             output reg [3:0] q);

    always @ (posedge clk)
        q <= d;
endmodule
```

→ register logic is performed using always statements (always @ (sensitivity list):)

→ signals keep their values unless an event in the sensitivity list occurs

→ usually it is a clock, as in this case → posedge: rising edge, negedge: falling edge

→ we don't use assign statements here, those are for CL blocks outside of always statements

→ instead, we use the nonblocking assignment (<=), which is used for seq. logic and runs concurrently (no order) unlike =

→ all signals on the LHS of always statements must be declared a register → example: output reg [3:0] q

example: resettable register, allowing for both sync. and async. reset
(2 modules)

```
module flopr (input clk,
              input reset,
              input [3:0] d,
              output reg [3:0] q);

    // asynchronous reset
    always @ (posedge clk, posedge reset)
        if (reset) q <= 4'b0;
        else q <= d;
endmodule

module flopr (input clk,
              input reset,
              input [3:0] d,
              output reg [3:0] q);

    // synchronous reset
    always @ (posedge clk)
        if (reset) q <= 4'b0;
        else q <= d;
endmodule
```

→ multiple elements in an always sens. list are separated by a comma or "or"

→ you will see, we have if statements, for loops, while loops ...

synchronizer:

```
module sync (input clk,
             input d,
             output reg q);

    reg n1;

    always @ (posedge clk)
        begin
            n1 <= d;
            q <= n1;
        end
endmodule
```



→ begin and end are necessary in any statement or block

that makes more than one assignment, good practice otherwise

Latch:

```
module latch (input clk,
              input [3:0] d,
              output reg [3:0] q);

    always @ (clk, d)
        if (clk) q <= d;
endmodule
```



→ always @ (clk, d) recomputes statements inside whenever clk or d changes

→ always @ (*) recomputes the statements inside whenever any signal on RHS of = or <= changes

→ blocking assignment (=) used for CL in always blocks, they run sequentially in-order

→ always blocks mostly for sequential blocks, but can behaviorally describe CL if the sens. list

responds to all input changes and the body gives outputs in all cases

→ the book calls always blocks "always/process statements"

→ case and if can only be used in always or initial blocks (initial does not synthesize and is used for testbenches)

```
module sevenseg (input [3:0] data,
                 output reg [6:0] segments);

    always @ (*)
        case (data)
            // abc_defg
            0: segments = 7'b111_1110;
            1: segments = 7'b011_0000;
            2: segments = 7'b110_1101;
            3: segments = 7'b111_1001;
            4: segments = 7'b011_0011;
            5: segments = 7'b101_1011;
            6: segments = 7'b101_1111;
            7: segments = 7'b111_0000;
            8: segments = 7'b111_1111;
            9: segments = 7'b111_1011;
            default: segments = 7'b000_0000;
        endcase
endmodule
```

→ case checks value of an input (here data) in base 10


→ if-else statements can be made like in regular programming, just follow the begin/end rules

→ case z includes don't cares from T1's as?

→ rule in always blocks: don't assign to the same signal more than once

```
module priority_casez (input [3:0] a,
                     output reg [3:0] y);

    always @ (*)
        casez (a)
            4'b111?: y = 4'b1000;
            4'b011?: y = 4'b0100;
            4'b001?: y = 4'b0010;
            4'b0001: y = 4'b0001;
            default: y = 4'b0000;
        endcase
endmodule
```



FSM:

Verilog (divides CLK by 3)

```
module divideby3FSM (input clk,
                    input reset,
                    output y);

    reg [1:0] state, nextstate;

    parameter S0 = 2'b00;
    parameter S1 = 2'b01;
    parameter S2 = 2'b10;

    // state register
    always @ (posedge clk, posedge reset)
        if (reset) state <= S0;
        else state <= nextstate;


    // next state logic
    always @ (*)
        case (state)
            S0: nextstate = S1;
            S1: nextstate = S2;
            S2: nextstate = S0;
            default: nextstate = S0;
        endcase

    // output logic
    assign y = (state == S0);
endmodule
```

→ parameter statements are used to define constants and are standard for any FSM state


→ default in case needed to fulfill our def. of combinational for next state logic

→ notice the equality comparison == or inequality != with strict versions === and !==



FSM Example 2: Recall the Smiling Snail

- Alyssa P. Hacker has a snail that crawls down a paper tape with 1's and 0's on it
- The snail smiles whenever the last four digits it has crawled over are 1101
- Design Moore and Mealy FSMs of the snail's brain



Verilog (Moore-style)

```
module patternMoore (input clk,
                    input reset,
                    input a,
                    output y);

    reg [2:0] state, nextstate;

    parameter S0 = 3'b000;
    parameter S1 = 3'b001;
    parameter S2 = 3'b010;
    parameter S3 = 3'b011;
    parameter S4 = 3'b100;

    // state register
    always @ (posedge clk, posedge reset)
        if (reset) state <= S0;
        else state <= nextstate;

    // next state logic
    always @ (*)
        case (state)
            S0: if (a) nextstate = S1;
                else nextstate = S0;
            S1: if (a) nextstate = S2;
                else nextstate = S0;
            S2: if (a) nextstate = S2;
                else nextstate = S3;
            S3: if (a) nextstate = S4;
                else nextstate = S0;
            S4: if (a) nextstate = S2;
                else nextstate = S0;
            default: nextstate = S0;
        endcase

    // output logic
    assign y = (state == S4);
endmodule
```

Verilog (Mealy style)

```
module patternMealy (input clk,
                    input reset,
                    input a,
                    output y);

    reg [1:0] state, nextstate;

    parameter S0 = 2'b00;
    parameter S1 = 2'b01;
    parameter S2 = 2'b10;
    parameter S3 = 2'b11;

    // state register
    always @ (posedge clk, posedge reset)
        if (reset) state <= S0;
        else state <= nextstate;

    // next state logic
    always @ (*)
        case (state)
            S0: if (a) nextstate = S1;
                else nextstate = S0;
            S1: if (a) nextstate = S2;
                else nextstate = S0;
            S2: if (a) nextstate = S2;
                else nextstate = S3;
            S3: if (a) nextstate = S1;
                else nextstate = S0;
            default: nextstate = S0;
        endcase


    // output logic
    assign y = (a & state == S3);
endmodule
```

Parameterizing:

Verilog

```
module mux2
    #(parameter width = 8)
    (input [width-1:0] d0, d1,
     input [width-1:0] s,
     output [width-1:0] y);
    assign y = s ? d1 : d0;
endmodule
```

default value
depends on width
we use a variable to decide module's width



```
module mux4_8 (input [7:0] d0, d1, d2, d3,
               input [1:0] s,
               output [7:0] y);

    wire [7:0] low, hi;

    mux2 lowmux(d0, d1, s[0], low);
    mux2 himux(d2, d3, s[1], hi);
    mux2 outmux(low, hi, s[1], y);
endmodule
```

→ you can instantiate like normal - 8 is default



```
module mux4_12 (input [11:0] d0, d1, d2, d3,
                input [1:0] s,
                output [11:0] y);

    wire [11:0] low, hi;

    mux2 #(12) lowmux(d0, d1, s[0], low);
    mux2 #(12) himux(d2, d3, s[1], hi);
    mux2 #(12) outmux(low, hi, s[1], y);
endmodule
```

→ need to instantiate with parameter 12

Testbenches: ~ HDL module used to test another module, called Device Under Test (DUT)

→ inputs and expected outputs → test vectors

Simple testbenches: instantiate DUT, apply inputs in initial block, use blocking assignments and delays, then the user manually verifies correct outputs, simulated like any module, not synthesizable

→ example: sillyfunction is a simple 3-input logic function → test vectors

```
module testbench1();
    reg a, b, c;
    wire y;

    // instantiate device under test
    sillyfunction dut(a, b, c, y);

    // apply inputs one at a time
    initial begin
        a = 0; b = 0; c = 0; #10;
        c = 1; #10;
        b = 1; c = 0; #10;
        c = 1; #10;
        a = 1; b = 0; c = 0; #10;
        c = 1; #10;
        b = 1; c = 0; #10;
        c = 1; #10;
    end
endmodule
```

→ initial statement: execute at start of simulation (0 time units) / here, we delay by 10 units per input

→ initial doesn't synthesize to hardware, signals assigned to in initial must be reg

self-checking testbench: more like a unit test, less tedious to verify → check against expected outputs, log failures

```
module testbench2();
  reg a, b, c;
  wire y;

  // instantiate device under test
  sillyfunction dut(a, b, c, y);

  // apply inputs one at a time
  // checking results
  initial begin
    a = 0; b = 0; c = 0; #10;
    if (y != 1) $display("000 failed.");
    c = 1; #10;
    if (y != 0) $display("001 failed.");
    b = 1; c = 0; #10;
    if (y != 0) $display("010 failed.");
    c = 1; #10;
    if (y != 0) $display("011 failed.");
    a = 1; b = 0; c = 0; #10;
    if (y != 1) $display("100 failed.");
    c = 1; #10;
    if (y != 1) $display("101 failed.");
    b = 1; c = 0; #10;
    if (y != 0) $display("110 failed.");
    c = 1; #10;
    if (y != 0) $display("111 failed.");
  end
endmodule
```

→ \$ is a prefix for system tasks / \$display prints a message to the simulator console

→ only strict equality works well with x and z

using a test vector file is more efficient for more vectors:

→ binary text file

→ we use a separate test clock to simulate until file is done

→ \$readmemb reads a binary file into the testvectors array,

```
module testbench3();
  reg clk, reset;
  reg a, b, c, yexpected;
  wire y;
  reg[31:0] vectornum, errors;
  reg[3:0] testvectors[10000:0];
  // instantiate device under test
  sillyfunction dut(a, b, c, y);

  // generate clock
  always
  begin
    clk = 1; #5; clk = 0; #5;
  end

  // at start of test, load vectors
  // and pulse reset
  initial
  begin
    $readmemb("example.tv", testvectors);
    vectornum = 0; errors = 0;
    reset = 1; #27; reset = 0;
  end

  // apply test vectors on rising edge of clk
  always @(posedge clk)
  begin
    #1; {a, b, c, yexpected} =
      testvectors[vectornum];
  end

  // check results on falling edge of clk
  always @(negedge clk)
  if (~reset) begin // skip during reset
    if (y != yexpected) begin
      $display("Error: inputs = %b", {a, b, c});
      $display(" outputs = %b (%b expected)",
        y, yexpected);
      errors = errors + 1;
    end
    vectornum = vectornum + 1;
    if (testvectors[vectornum] == 4'bxx) begin
      $display("%d tests completed with %d errors",
        vectornum, errors);
    end
  end
endmodule
```

\$readmemb does the same for hex files

→ wait one unit of time after rising edge to set new vector and expected output using bit swizzling

→ on falling edge, check if reset not active, error messages if output doesn't match expectation


→ \$display has this syntax foo: \$display("custom message %0b", {a,b,c})

→ %0b prints the incoming number in binary, %d in decimal

→ \$finish terminates simulation

Automatic Testbenches:

Golden Models: ideal circuit, may be written in other language



→ we have to automatically generate test vectors → sequential coverage or randomization?

pseudocode →

```
module testbench1();
  ... // variable declarations, clock, etc.

  // instantiate device under test
  sillyfunction dut(a, b, c, y_dut);
  golden_model gold(a, b, c, y_gold);

  // instantiate test pattern generator
  test_pattern_generator tgen(a, b, c, clk);

  // check if y_dut is ever not equal to y_gold
  always @(negedge clk)
  begin
    if (y_dut != y_gold)
      $display(...);
  end
endmodule
```

→ this is fully automated / scalable, even timing can be compared with a golden timing model

→ challenges: creating a golden model, test generation, brute forcing tests is inefficient

Van Neumann Model, Processor Design:


1/0n Neumann Model:

(smallest pieces of work)

→ a program consists of instructions and is contained in memory, the control unit orders instructions

Memory: if a memory has 2^8 distinct locations and each¹ can store 8 bits

it has an address space of 2^8 locations and addressability of 8 bits → byte addressable, some ISA's word-addressable



→ reading it: place address in MAR (Memory Access Register) → "interrogate memory"

→ the information stored in that MAR location is then placed in the MDR (Memory Data Register)

→ writing to mem.: put the address in MAR, value in MDR

→ interrogate mem. with $WE=1$ (write enable) → MDR info written to the location we wrote into the MAR

→ MAR locations are permanent, what's inside them isn't

000	
001	
010	
011	
100	00001110
101	
110	00000100
111	

Processing Unit: (PU)

→ the simplest PU is an ALU with ADD, SUBTRACT, basic logic

→ the size of quantities processed in bits is the word length, the element is called a word (RF)

→ a register file stores results temporarily (near the ALU) → LC-3 has 8 16-bit registers for this

I/O devices: peripherals like keyboards and displays

Control Unit: contains Instruction Register (IR) that contains the current ins., then a register that contains the next ins.'s address (program counter, abbreviated to PC)

LC-3 as a Von Neumann machine:

→ filled-in arrowheads → data flows using the signal/bus

→ not filled-in → control signals, no data flows

Memory: our MAR (because of an address space of 2^{16} locations)

and MDR (because of 16-bit words) are 16-bit

I/O: the KBDR (keyboard data reg.) contains the ASCII of struck keys,

the KBSR contains their status

→ the monitor likewise has a Display Data Register and Display Status Reg.

For the same purposes

Processing Unit: consists of ALU and RF

Control Unit: FSM directs activity, processing is done cycle-by-cycle (CLK),

IR is an input for the FSM, the PC saves the next ins.

→ The FSM outputs the 2-bit opcode ALUK, and control signals like GateALU (determines if ALU output pushed onto processor bus)

Instruction Processing:

→ an instruction consists of an opcode (describing what's done) and operands (where it is done)




Figure 4.3 The LC-3 as an example of the von Neumann model